

```
# importing required packages

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings(action = 'ignore')

def heading(info):
    print("\n\n##### {} #####".format(info))

# read the dataset
dataSet = pd.read_csv('Bangalore_traffic_Dataset.csv', encoding = 'unicode_escape')
```

```
# print info about the data
dataSet.info()
heading("Sample data points from the dataset")
dataSet.head(5)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8936 entries, 0 to 8935
Data columns (total 16 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Date                                     8936 non-null   object
1   Area Name                               8936 non-null   object
2   Road/Intersection Name                  8936 non-null   object
3   Traffic Volume                          8936 non-null   int64
4   Average Speed                           8936 non-null   float64
5   Travel Time Index                       8936 non-null   float64
6   Congestion Level                        8936 non-null   float64
7   Road Capacity Utilization               8936 non-null   float64
8   Incident Reports                        8936 non-null   int64
9   Environmental Impact                    8936 non-null   float64
10  Public Transport Usage                   8936 non-null   float64
11  Traffic Signal Compliance                8936 non-null   float64
12  Parking Usage                           8936 non-null   float64
13  Pedestrian and Cyclist Count             8936 non-null   int64
14  Weather Conditions                       8936 non-null   object
15  Roadwork and Construction Activity       8936 non-null   object
dtypes: float64(8), int64(3), object(5)
memory usage: 1.1+ MB
```

```
##### Sample data points from the dataset #####
```

	Date	Area Name	Road/Intersection Name	Traffic Volume	Average Speed	Travel Time Index	Congestion Level
0	2022-01-01	Indiranagar	100 Feet Road	50590	50.230299	1.500000	100.000000
1	2022-01-01	Indiranagar	CMH Road	30825	29.377125	1.500000	100.000000
2	2022-01-01	Whitefield	Marathahalli Bridge	7399	54.474398	1.039069	28.347990
3	2022-01-01	Koramangala	Sony World Junction	60874	43.817610	1.500000	100.000000
4	2022-01-01	Koramangala	Sarjapur Road	57292	41.116763	1.500000	100.000000

```
# lets find the individual column statistics
heading("Stats about non-numeric values")
print(dataSet.describe(include = "object"))

heading("Stats about numeric values")
print(dataSet.describe(include = "number"))
```



```
##### Stats about non-numeric values #####
```

	Date	Area Name	Road/Intersection	Name	Weather	Conditions
count	8936	8936		8936		8936
unique	952	8		16		5
top	2023-01-24	Indiranagar	100 Feet Road			Clear
freq	15	1720		860		5426

```
Roadwork and Construction Activity
```

count	8936
unique	2
top	No
freq	8054

```
##### Stats about numeric values #####
```

	Traffic Volume	Average Speed	Travel Time Index	Congestion Level
count	8936.000000	8936.000000	8936.000000	8936.000000
mean	29236.048120	39.447427	1.375554	80.818041
std	13001.808801	10.707244	0.165319	23.533182
min	4233.000000	20.000000	1.000039	5.160279
25%	19413.000000	31.775825	1.242459	64.292905
50%	27600.000000	39.199368	1.500000	92.389018
75%	38058.500000	46.644517	1.500000	100.000000
max	72039.000000	89.790843	1.500000	100.000000

	Road Capacity Utilization	Incident Reports	Environmental Impact \
count	8936.000000	8936.000000	8936.000000
mean	92.029215	1.570389	108.472096
std	16.583341	1.420047	26.003618
min	18.739771	0.000000	58.466000
25%	97.354990	0.000000	88.826000
50%	100.000000	1.000000	105.200000
75%	100.000000	2.000000	126.117000
max	100.000000	10.000000	194.078000

	Public Transport Usage	Traffic Signal Compliance	Parking Usage \
count	8936.000000	8936.000000	8936.000000
mean	45.086651	79.950243	75.155597
std	20.208460	11.585006	14.409394
min	10.006853	60.003933	50.020411
25%	27.341191	69.828270	62.545895
50%	45.170684	79.992773	75.317610
75%	62.426485	89.957358	87.518589


max 79.979744 99.993652 99.995049

```
Pedestrian and Cyclist Count
count      8936.000000
mean       114.533348
std        36.812573
min        66.000000
25%        94.000000
50%       102.000000
75%       111.000000
max       243.000000
```

```
# lets first verify how to target variable is distributed
heading("Target variable \"Traffic volume\" distribution")
sns.histplot(data = dataSet, x = "Traffic Volume")
```



```
##### Target variable "Traffic volume" distribution #####
<Axes: xlabel='Traffic Volume', ylabel='Count'>
```



```
# lets convert the categorical values to numeric
def convert_categorical_to_numeric(dataframe, categorical_cols):

    for col in categorical_cols:
        if col in dataframe.columns:
            # create a mapping for the unique values in the column
            unique_values = dataframe[col].unique()
            value_mapping = {label: idx for idx, label in enumerate(unique_val

            # apply the mapping to convert to numeric
            dataframe[col] = [value_mapping[val] for val in dataframe[col]]

    return dataframe

# leaving date column as of now and converting other columns

# we will backup the original dataset
originalDataset = dataSet.copy()

# select the relevant columns and convert them
columnsToConvert = ["Roadwork and Construction Activity", "Weather Conditions", "
dataSet = convert_categorical_to_numeric(dataSet, columnsToConvert)

heading("After conversion to numeric values")
print(dataSet[columnsToConvert].head())
```



After conversion to numeric values

	Roadwork and Construction Activity	Weather Conditions	Area Name \
0	0	0	0
1	0	0	0
2	0	0	1
3	0	0	2
4	0	0	2

	Road/Intersection Name
0	0
1	1
2	2
3	3
4	4

```
# drop unrequired columns based on corelation matrix
dropThem = ["Public Transport Usage", "Traffic Signal Compliance", "Parking Usage"]
dataSet = dataSet.drop(columns=dropThem)
```

```
heading("Final dataset columns")
print(dataSet.head())
```



Final dataset columns

	Area Name	Road/Intersection Name	Traffic Volume	Average Speed	\
0	0	0	50590	50.230299	
1	0	1	30825	29.377125	
2	1	2	7399	54.474398	
3	2	3	60874	43.817610	
4	2	4	57292	41.116763	

	Travel Time Index	Congestion Level	Road Capacity Utilization	\
0	1.500000	100.000000	100.000000	
1	1.500000	100.000000	100.000000	
2	1.039069	28.347994	36.396525	
3	1.500000	100.000000	100.000000	
4	1.500000	100.000000	100.000000	

	Incident Reports	Environmental Impact	Pedestrian and Cyclist Count
0	0	151.180	111
1	1	111.650	100
2	0	64.798	189
3	1	171.748	111
4	3	164.584	104

```
# seperate the input and target columns into numpy arrays
if isinstance(dataSet, pd.DataFrame):
    dataSet = dataSet.to_numpy()
# print(dataset)
X = dataSet[:, [0, 1, 3, 4,5,6,7,8,9]]
Y = dataSet[:, 2]

# adding extra column for intercepts
X = np.hstack((np.ones((X.shape[0], 1)), X))
heading("Printing X and Y variables for the model")
print(X[:5])
print(Y[:5])
```



```
##### Printing X and Y variables for the model #####
[[ 1.         0.         0.         50.23029856  1.5
   100.        100.        0.         151.18        111.         ]
 [ 1.         0.         1.         29.37712471  1.5
   100.        100.        1.         111.65        100.         ]
 [ 1.         1.         2.         54.47439821  1.03906885
  28.34799386  36.39652494  0.         64.798        189.         ]
 [ 1.         2.         3.         43.81761039  1.5
   100.        100.        1.         171.748        111.         ]
 [ 1.         2.         4.         41.11676289  1.5
   100.        100.        3.         164.584        104.         ]]
[50590. 30825.  7399. 60874. 57292.]
```

```
# shuffle the datasets
indices = np.arange(X.shape[0])
np.random.shuffle(indices)
#
X_shuffled = X[indices]
Y_shuffled = Y[indices]

# split the dataset into 80:20
split_ratio = 0.2
split_index = int(len(X_shuffled) * split_ratio)

X_train = X_shuffled[:split_index]
Y_train = Y_shuffled[:split_index]

X_test = X_shuffled[split_index:]
Y_test = Y_shuffled[split_index:]

print("Training set samples: ", X_train.shape[0])
print("Testing set samples: ", X_test.shape[0])

↔ Training set samples: 1787
   Testing set samples: 7149
```



```
# for computing gradient descent, we use the the firmula new_weights = old_weig
#  $dJ/dW = -2X^T Y + 2X^T XW = 2X^T(XW-Y)$ 
```

```
def gradient_descent(X, Y, learning_rate=0.01, iterations=1000):
    n_samples, n_features = X.shape

    # initialize weights to 0
    weights = np.zeros(n_features)
    # List to store cost at each iteration for plotting convergence
    cost_history = []

    for i in range(iterations):
        # Y_pred -> XW
        Y_pred = X @ weights

        # compute the error -> XW-Y
        error = Y_pred - Y

        # compute gradient -> dJ/dW
        gradient = (2 / n_samples) * (X.T @ error)

        # Update weights
        weights -= learning_rate * gradient

        # Compute Mean Squared Error (Cost Function)
        cost = (1 / n_samples) * np.sum(error ** 2)

        # (Optional) Print cost at intervals
        if i % 50 == 0:
            print(f"Iteration {i+1}: Cost {cost}")
            cost_history.append(cost)

    return weights, cost_history
```


```
# we solve the least squares problem using the gradient descent algorithm
weights, cost_history = gradient_descent(X_train,Y_train,0.00002,1500)
heading("Training on linear regression with given dataset for 10 iterations")
print(weights[:5])
```

```
↪ Iteration 1: Cost 1024440498.4297707
Iteration 51: Cost 19257093.994292814
Iteration 101: Cost 10232535.491425855
Iteration 151: Cost 6051920.219076406
Iteration 201: Cost 4053860.8857249906
Iteration 251: Cost 3060828.882377268
Iteration 301: Cost 2543862.3888015794
Iteration 351: Cost 2260861.998928928
Iteration 401: Cost 2098092.154895537
Iteration 451: Cost 2000241.9044550362
Iteration 501: Cost 1939238.3680511315
Iteration 551: Cost 1900124.1144997755
Iteration 601: Cost 1874519.4331545122
Iteration 651: Cost 1857503.5425434855
Iteration 701: Cost 1846068.2464918403
Iteration 751: Cost 1838315.207653634
Iteration 801: Cost 1833017.9106328539
Iteration 851: Cost 1829370.567145155
Iteration 901: Cost 1826837.6907410177
Iteration 951: Cost 1825060.655838286
Iteration 1001: Cost 1823798.0246972395
Iteration 1051: Cost 1822886.668808727
Iteration 1101: Cost 1822216.0840303628
Iteration 1151: Cost 1821711.262183427
Iteration 1201: Cost 1821321.1974394964
Iteration 1251: Cost 1821011.142274124
Iteration 1301: Cost 1820757.3758214433
Iteration 1351: Cost 1820543.6633198147
Iteration 1401: Cost 1820358.8571920535
Iteration 1451: Cost 1820195.2703071258
```

```
##### Training on linear regression with given dataset for 10 iterations ##
[ -5.4905183 -21.597431 -56.54143836 -35.13642342 -5.22527202]
```

```
# Plotting the cost history using seaborn
print(len(cost_history))
plt.figure(figsize=(10,6))
sns.lineplot(x=range(len(cost_history)), y=cost_history, color='blue')
plt.xlabel('Iterations')
plt.ylabel('Cost (MSE)')
plt.title('Cost Reduction over Iterations')
plt.show()
```

↩ 30



```
# lets predict values
Y_pred = X_test @ weights

# we will use mean absolute percentage error to calculate the error percentage
MAPE = np.mean(np.abs((Y_test - Y_pred) / Y_test)) * 100

heading("Printing the MAPE and first 10 predictions with actual values")
print("MAPE: {} %".format(MAPE))
for i in range(10):
    print("\nPredicted value: {0} \t Actual value: {1}".format(Y_pred[i], Y_test[i]))
```



```
##### Printing the MAPE and first 10 predictions with actual values #####
MAPE: 5.557551943610017 %
```

Predicted value: 34426.65982240805	Actual value: 35269.0
Predicted value: 27926.53382259589	Actual value: 28038.0
Predicted value: 27379.23702980667	Actual value: 27049.0
Predicted value: 23692.006650849784	Actual value: 24265.0
Predicted value: 40315.83882924316	Actual value: 41557.0
Predicted value: 10393.53185509519	Actual value: 8059.0
Predicted value: 9942.332775055342	Actual value: 8077.0
Predicted value: 11039.920832360887	Actual value: 14725.0
Predicted value: 11215.776744104967	Actual value: 12832.0
Predicted value: 35995.83136993581	Actual value: 36190.0

