

```
# importing required packages

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings(action = 'ignore')

def heading(info):
    print("\n\n##### {} #####".format(info))

# read the dataset
dataSet = pd.read_csv('Bangalore_traffic_Dataset.csv', encoding = 'unicode_escape')
```

```
# print info about the data
dataSet.info()
heading("Sample data points from the dataset")
dataSet.head(5)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8936 entries, 0 to 8935
Data columns (total 16 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Date                                       8936 non-null   object
1   Area Name                                8936 non-null   object
2   Road/Intersection Name                    8936 non-null   object
3   Traffic Volume                            8936 non-null   int64
4   Average Speed                            8936 non-null   float64
5   Travel Time Index                        8936 non-null   float64
6   Congestion Level                         8936 non-null   float64
7   Road Capacity Utilization                8936 non-null   float64
8   Incident Reports                         8936 non-null   int64
9   Environmental Impact                    8936 non-null   float64
10  Public Transport Usage                   8936 non-null   float64
11  Traffic Signal Compliance                8936 non-null   float64
12  Parking Usage                           8936 non-null   float64
13  Pedestrian and Cyclist Count            8936 non-null   int64
14  Weather Conditions                      8936 non-null   object
15  Roadwork and Construction Activity       8936 non-null   object
dtypes: float64(8), int64(3), object(5)
memory usage: 1.1+ MB
```

```
##### Sample data points from the dataset #####
```

	Date	Area Name	Road/Intersection Name	Traffic Volume	Average Speed	Travel Time Index	Congestion Level
0	2022-01-01	Indiranagar	100 Feet Road	50590	50.230299	1.500000	100.000000
1	2022-01-01	Indiranagar	CMH Road	30825	29.377125	1.500000	100.000000
2	2022-01-01	Whitefield	Marathahalli Bridge	7399	54.474398	1.039069	28.347990
3	2022-01-01	Koramangala	Sony World Junction	60874	43.817610	1.500000	100.000000
4	2022-01-01	Koramangala	Sarjapur Road	57292	41.116763	1.500000	100.000000

```
# lets find the individual column statistics
heading("Stats about non-numeric values")
print(dataSet.describe(include = "object"))

heading("Stats about numeric values")
print(dataSet.describe(include = "number"))
```



Stats about non-numeric values

	Date	Area Name	Road/Intersection	Name	Weather	Conditions
count	8936	8936		8936		8936
unique	952	8		16		5
top	2023-01-24	Indiranagar	100 Feet Road			Clear
freq	15	1720		860		5426

Roadwork and Construction Activity

count	8936
unique	2
top	No
freq	8054

Stats about numeric values

	Traffic Volume	Average Speed	Travel Time Index	Congestion Level
count	8936.000000	8936.000000	8936.000000	8936.000000
mean	29236.048120	39.447427	1.375554	80.818041
std	13001.808801	10.707244	0.165319	23.533182
min	4233.000000	20.000000	1.000039	5.160279
25%	19413.000000	31.775825	1.242459	64.292905
50%	27600.000000	39.199368	1.500000	92.389018
75%	38058.500000	46.644517	1.500000	100.000000
max	72039.000000	89.790843	1.500000	100.000000

	Road Capacity Utilization	Incident Reports	Environmental Impact \
count	8936.000000	8936.000000	8936.000000
mean	92.029215	1.570389	108.472096
std	16.583341	1.420047	26.003618
min	18.739771	0.000000	58.466000
25%	97.354990	0.000000	88.826000
50%	100.000000	1.000000	105.200000
75%	100.000000	2.000000	126.117000
max	100.000000	10.000000	194.078000

	Public Transport Usage	Traffic Signal Compliance	Parking Usage \
count	8936.000000	8936.000000	8936.000000
mean	45.086651	79.950243	75.155597
std	20.208460	11.585006	14.409394
min	10.006853	60.003933	50.020411
25%	27.341191	69.828270	62.545895
50%	45.170684	79.992773	75.317610
75%	62.426485	89.957358	87.518589


max 79.979744 99.993652 99.995049

```
Pedestrian and Cyclist Count
count      8936.000000
mean       114.533348
std        36.812573
min        66.000000
25%        94.000000
50%       102.000000
75%       111.000000
max       243.000000
```

```
# lets first verify how to target variable is distributed
heading("Target variable \"Traffic volume\" distribution")
sns.histplot(data = dataSet, x = "Traffic Volume")
```



```
##### Target variable "Traffic volume" distribution #####
<Axes: xlabel='Traffic Volume', ylabel='Count'>
```



```
# lets convert the categorical values to numeric
def convert_categorical_to_numeric(dataframe, categorical_cols):

    for col in categorical_cols:
        if col in dataframe.columns:
            # create a mapping for the unique values in the column
            unique_values = dataframe[col].unique()
            value_mapping = {label: idx for idx, label in enumerate(unique_valu

            # apply the mapping to convert to numeric
            dataframe[col] = [value_mapping[val] for val in dataframe[col]]

    return dataframe

# leaving date column as of now and converting other columns

# we will backup the original dataset
originalDataset = dataSet.copy()

# select the relevant columns and convert them
columnsToConvert = ["Roadwork and Construction Activity", "Weather Conditions", "
dataSet = convert_categorical_to_numeric(dataSet, columnsToConvert)

heading("After conversion to numeric values")
print(dataSet[columnsToConvert].head())
```



```
##### After conversion to numeric values #####
```

	Roadwork and Construction Activity	Weather Conditions	Area Name \
0	0	0	0
1	0	0	0
2	0	0	1
3	0	0	2
4	0	0	2

	Road/Intersection Name
0	0
1	1
2	2
3	3
4	4

```
# drop unrequired columns based on corelation matrix
dropThem = ["Public Transport Usage", "Traffic Signal Compliance", "Parking Usage"]
dataSet = dataSet.drop(columns=dropThem)
```

```
heading("Final dataset columns")
print(dataSet.head())
```



Final dataset columns

	Area Name	Road/Intersection Name	Traffic Volume	Average Speed	\
0	0	0	50590	50.230299	
1	0	1	30825	29.377125	
2	1	2	7399	54.474398	
3	2	3	60874	43.817610	
4	2	4	57292	41.116763	

	Travel Time Index	Congestion Level	Road Capacity Utilization	\
0	1.500000	100.000000	100.000000	
1	1.500000	100.000000	100.000000	
2	1.039069	28.347994	36.396525	
3	1.500000	100.000000	100.000000	
4	1.500000	100.000000	100.000000	

	Incident Reports	Environmental Impact	Pedestrian and Cyclist Count
0	0	151.180	111
1	1	111.650	100
2	0	64.798	189
3	1	171.748	111
4	3	164.584	104

```
# seperate the input and target columns into numpy arrays
if isinstance(dataSet, pd.DataFrame):
    dataSet = dataSet.to_numpy()
# print(dataset)
X = dataSet[:, [0, 1, 3, 4,5,6,7,8,9]]
Y = dataSet[:, 2]

# adding extra column for intercepts
X = np.hstack((np.ones((X.shape[0], 1)), X))
heading("Printing X and Y variables for the model")
print(X[:5])
print(Y[:5])
```



```
##### Printing X and Y variables for the model #####
[[ 1.         0.         0.         50.23029856  1.5
   100.        100.        0.         151.18        111.         ]
 [ 1.         0.         1.         29.37712471  1.5
   100.        100.        1.         111.65         100.         ]
 [ 1.         1.         2.         54.47439821  1.03906885
   28.34799386  36.39652494  0.         64.798         189.         ]
 [ 1.         2.         3.         43.81761039  1.5
   100.        100.        1.         171.748         111.         ]
 [ 1.         2.         4.         41.11676289  1.5
   100.        100.        3.         164.584         104.         ]]
[50590. 30825.  7399. 60874. 57292.]
```

```
# shuffle the datasets
indices = np.arange(X.shape[0])
np.random.shuffle(indices)
#
X_shuffled = X[indices]
Y_shuffled = Y[indices]
```

```
# we normalize our dataset

# we use standadization for input variables, since they are of different scales
# standardization helps set the mean to 0 and std dev to 1.
intercept_column = X_shuffled[:, 0]
features = X_shuffled[:, 1:]

X_mean = np.mean(features, axis=0)
X_std = np.std(features, axis=0)

# Normalize features
X_norm_temp = (features - X_mean) / X_std
X_norm = np.column_stack((intercept_column, X_norm_temp))

# we use min max normalization for the output variable, since it is in fixed range
# also, standardization may yeild negative values which may not be as intuitive
Y_mean = np.mean(Y, axis=0)
Y_std = np.std(Y, axis=0)

# Normalize features
Y_norm = (Y - Y_mean) / Y_std
```



```
# split the dataset into 80:20
split_ratio = 0.8
split_index = int(len(X_shuffled) * split_ratio)

X_train = X_norm[:split_index]
Y_train = Y_norm[:split_index]

X_test = X_norm[split_index:]
Y_test = Y_norm[split_index:]

print("Training set samples size: ", X_train.shape[0])
print("Testing set samples size: ", X_test.shape[0])

print("Training set samples: ", X_train[:5])
print("Testing set samples: ", Y_train[:5])
```

```
↔ Training set samples size: 7148
Testing set samples size: 1788
Training set samples: [[ 1.          0.04836864 -0.26557885 -0.24502419  0
  0.48067702 -1.10593338  0.98437596 -0.39481515]
 [ 1.          0.51563584  0.42430854  0.61332466 -0.82137174 -2.20155768
 -3.06246808 -1.10593338 -1.58425655  2.21313545]
 [ 1.          0.04836864 -0.26557885  0.38672395 -2.10567337  0.17118203
  0.48067702 -1.10593338 -0.3993551  0.23000635]
 [ 1.         -1.35343297 -1.41539116 -1.81638862  0.75280715  0.81514828
  0.48067702  0.30254958  0.48295613 -0.17748593]
 [ 1.          0.04836864 -0.26557885  1.95049546  0.13711301 -1.99925888
 -1.5518211  -1.10593338 -1.28358925  1.85997547]]
Testing set samples: [ 1.64247507  0.1222169 -1.67963323  2.4334862  2.1
```

```
# for computing gradient descent, we use the the firmula new_weights = old_weig
#  $dJ/dW = -2X^T Y + 2X^T XW = 2X^T(XW-Y)$ 
```

```
def gradient_descent(X, Y, learning_rate=0.01, iterations=1000):
    n_samples, n_features = X.shape

    # initialize weights to 0
    weights = np.zeros(n_features)
    # List to store cost at each iteration for plotting convergence
    cost_history = []

    for i in range(iterations):
        # Y_pred -> XW
        Y_pred = X @ weights

        # compute the error -> XW-Y
        error = Y_pred - Y

        # compute gradient -> dJ/dW
        gradient = (2 / n_samples) * (X.T @ error)

        # Update weights
        weights -= learning_rate * gradient

        # Compute Mean Squared Error (Cost Function)
        cost = (1 / n_samples) * np.sum(error ** 2)

        # (Optional) Print cost at intervals
        if i % 100 == 0:
            print(f"Iteration {i+1}: Cost {cost}")
            cost_history.append(cost)

    return weights, cost_history
```


```
# we solve the least squares problem using the gradient descent algorithm
weights, cost_history = gradient_descent(X_train,Y_train,0.1,3000)
heading("Training on linear regression with given dataset for 10 iterations")
print(weights[:5])
```

```
↪ Iteration 1: Cost 1.0058789972297144
Iteration 101: Cost 1.0043202614022586
Iteration 201: Cost 1.0042667705344606
Iteration 301: Cost 1.0042605008369565
Iteration 401: Cost 1.0042597657038155
Iteration 501: Cost 1.0042596795076422
Iteration 601: Cost 1.004259669400927
Iteration 701: Cost 1.0042596682158895
Iteration 801: Cost 1.0042596680769411
Iteration 901: Cost 1.0042596680606488
Iteration 1001: Cost 1.0042596680587388
Iteration 1101: Cost 1.0042596680585145
Iteration 1201: Cost 1.0042596680584883
Iteration 1301: Cost 1.0042596680584854
Iteration 1401: Cost 1.0042596680584848
Iteration 1501: Cost 1.004259668058485
Iteration 1601: Cost 1.0042596680584848
Iteration 1701: Cost 1.0042596680584848
Iteration 1801: Cost 1.0042596680584848
Iteration 1901: Cost 1.0042596680584848
Iteration 2001: Cost 1.0042596680584848
Iteration 2101: Cost 1.0042596680584848
Iteration 2201: Cost 1.004259668058485
Iteration 2301: Cost 1.0042596680584848
Iteration 2401: Cost 1.0042596680584848
Iteration 2501: Cost 1.0042596680584848
Iteration 2601: Cost 1.0042596680584848
Iteration 2701: Cost 1.0042596680584848
Iteration 2801: Cost 1.0042596680584848
Iteration 2901: Cost 1.0042596680584848
```

```
##### Training on linear regression with given dataset for 10 iterations ##
[-0.00058856 -0.01247141  0.00304069 -0.00042342  0.00373874]
```

```
# Plotting the cost history using seaborn
print(len(cost_history))
plt.figure(figsize=(10,6))
sns.lineplot(x=range(len(cost_history)), y=cost_history, color='blue')
plt.xlabel('Iterations')
plt.ylabel('Cost (MSE)')
plt.title('Cost Reduction over Iterations')
plt.show()
```

↩ 30



```
# lets predict values
Y_pred = X_test @ weights
# we denormalize the predictions and actual testing data
Y_pred_original = (Y_pred * Y_std) + Y_mean
Y_test_denorm = (Y_test * Y_std) + Y_mean

# we will use mean absolute percentage error to calculate the error percentage
MAPE = np.mean(np.abs((Y_test_denorm - Y_pred_original) / Y_test_denorm)) * 100

heading("Printing the MAPE and first 10 predictions with actual values")
print("MAPE: {} %".format(MAPE))
for i in range(10):
    print("\nPredicted value: {0} \t Actual value: {1}".format(Y_pred_original[i], Y_test_denorm[i]))
```



```
##### Printing the MAPE and first 10 predictions with actual values #####
MAPE: 51.3920872695371 %
```

Predicted value: 29112.69893525145	Actual value: 51665.0
Predicted value: 28845.33853487075	Actual value: 36384.0
Predicted value: 29363.322605982026	Actual value: 41692.0
Predicted value: 29185.897419801313	Actual value: 29831.0
Predicted value: 29258.280767655135	Actual value: 10915.0
Predicted value: 29134.03138423169	Actual value: 25909.0
Predicted value: 29060.315434471842	Actual value: 37720.0
Predicted value: 29040.760707823174	Actual value: 26712.0
Predicted value: 28838.098747263044	Actual value: 33467.0
Predicted value: 29915.88538897134	Actual value: 27426.0

